

1. System overview

The ingestion stage accepts born-digital PDFs and raster scans. A router inspects the file signature and dispatches either the native text path or the OCR path.

2. Layout analysis

Region detection groups paragraphs, titles, tables and figures. Reading order is derived per column before fusion.

3. Table extraction

Structure recovery preserves merged cells and header relationships. Special glyphs survive only when the OCR charset covers them.

4. Diagnostics

Per-region confidence feeds the quality layer; failures are graded with explicit error codes for replay.

Symbol	Name	Unicode	Pipeline status
Glyph survival	matrix (merged header)	see rows	
✓	check	U+2713	expected: pass
⊗	tensor/cross	U+2297	expected: at risk
●	filled circle	U+25CF	expected: pass
○	open circle	U+25CB	expected: at risk

Footnote 1 — this document is generated deterministically by scripts/trial/generate_trial_samples.py. Any extraction regression is a pipeline change, not a sample drift.